



Waterwheel

Automated QA that runs inside **your own stack**

Waterwheel turns plain-English descriptions into real end-to-end browser tests — running in your environment, in your pipeline, with your choice of AI.

No brittle scripts. No data leaving your network.

Plain-English tests

Runs in your environment

Fits your CI & coding agents

Data never leaves

Every release runs the same slow, brittle gauntlet

Before a feature ships, your team has to prove the app still works. Manual passes are slow, and automated scripts break constantly — swallowing 40–50% of QA effort in upkeep alone.

🕒 Time per full test round — ~7 hours by hand



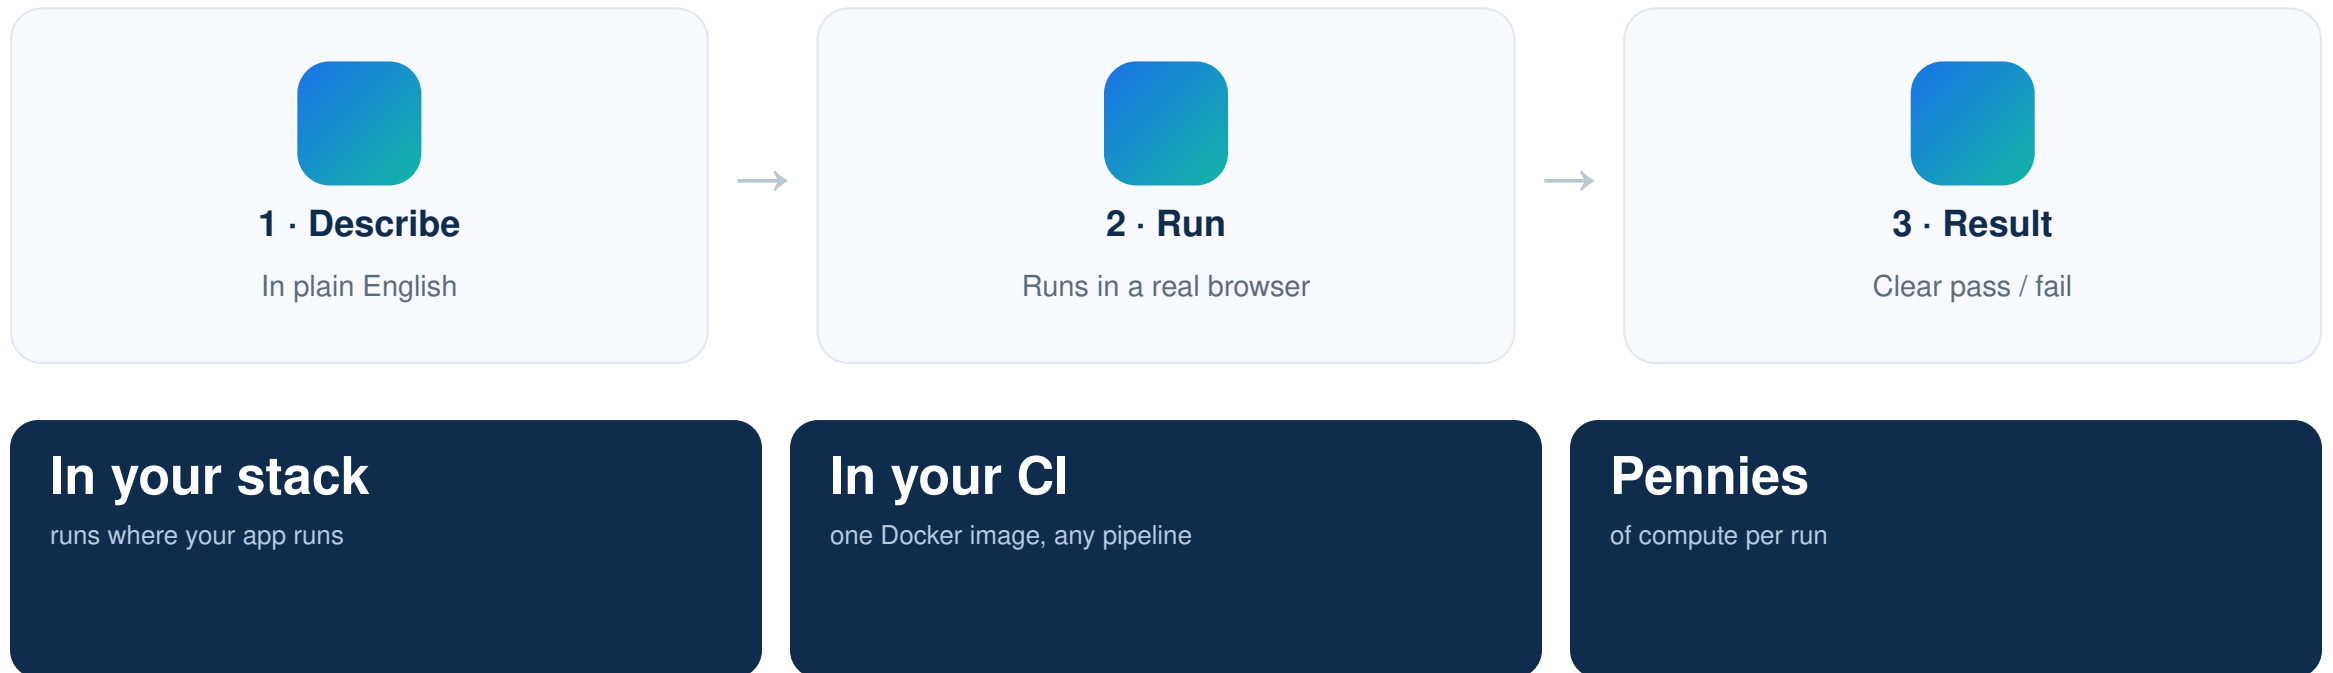
💰 Cost per full test round — ~\$390 by hand



Developer time Manual testing Writing test scripts Re-checking everything

Describe a test in plain English — Waterwheel runs it

It drives a real browser and reasons about the page, like your best manual tester. No selectors, no scripts to maintain.



Modern AI testing — without handing over your data

Complex, dependent flows — **sign up** → **activate by email** → **start a paid subscription** — handled end to end, inside your perimeter.

Capability	Script tools Selenium · Cypress	AI cloud tools testRigor · mabl	Waterwheel
Tests written in plain English	✗	✓	✓
Survives redesigns — no brittle selectors to fix	✗	✓	✓
Passes state between steps (generated logins, activation links)	✗	~	✓
Private test-and-fix loop with your coding agents	✗	✗	✓
Runs in your environment — data never leaves	✓	✗	✓
Your tests, your infra — no vendor lock-in	✓	✗	✓
Your choice of AI model (incl. self-hosted)	—	✗	✓

✓ Full support ~ Partial / limited ✗ Not supported

Where the leading AI test platforms fall short

The same gap runs through testRigor, mabl, TestSprite, QA.tech, and Testim: they're cloud SaaS you rent and don't control. Waterwheel closes each one.

THEIR LIMITATION

THE WATERWHEEL ANSWER

Cloud-only — your app, source, and test data run on their servers



Self-hosted in your own Docker — nothing leaves your network

Locked in the platform — tests live in their cloud, often no exportable code



You own the tests — plain-English Markdown files kept in your repo

Rented per seat & per test run — ~\$450–\$1,000+/mo, up to \$40k+/yr



Self-hosted licensing — no per-seat fees or per-test-run cloud metering

Their fixed AI model, running in their cloud



Bring your own model — any provider, or self-hosted so LLM traffic stays in-house

Opaque, limited control — reusable logic and self-healing you can't inspect



Full control — teach the agent with your own readable skills, swappable per environment

Drops into the pipeline you already run

No platform to adopt, no data to hand over — Waterwheel meets your engineering workflow where it is.

One Docker image

Pull it and run — locally, on a CI runner, or in your own cluster. Nothing to stand up, works in any pipeline.

A private test-and-fix loop

Connect your coding agents (Claude Code, Cursor) so they run tests, read the failures, and fix their own code — entirely inside your network, never a vendor cloud.

Bring your own AI

Any major provider, or a self-hosted model. Your key, your model choice, your cost control — no lock-in to ours.

Inside your perimeter

Source, app, and test data stay on your systems. Built for teams with real security and compliance requirements.

What your engineering org gets

Faster releases, far less test upkeep, and a QA layer that satisfies security and compliance.



~8 min

release-ready run, unattended — vs ~7 hrs by hand



~50%

less time lost to test maintenance



0

lines of test code to write



100%

inside your perimeter — compliance-ready



See it run in **your** **environment**

Pull the Docker image, point it at your app, and watch Waterwheel test a real user flow — in minutes, on your own infrastructure.

1 · Pull the image

2 · Point it at your app

3 · First test in minutes